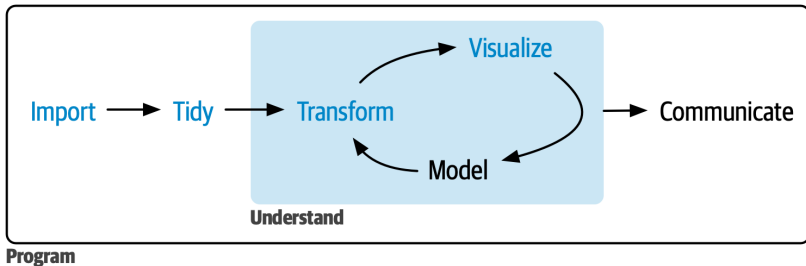


<https://r4ds.hadley.nz/> Chapter 3-9

CSCI 297b, Spring 2023

May 4, 2023

The Big Picture



The dplyr package and the nycflights13 dataset

```
library(nycflights13)  
library(tidyverse)
```

the nycflights13 dataset

```
> glimpse(flights)
Rows: 336,776
Columns: 19
$ year      <int> 2013, 2013, 2013, 2013,...
$ month     <int> 1, 1, 1, 1, 1, 1, 1, 1,...
$ day       <int> 1, 1, 1, 1, 1, 1, 1, 1,...
$ dep_time  <int> 517, 533, 542, 544, 554...
$ sched_dep_time <int> 515, 529, 540, 545, 600...
$ dep_delay <dbl> 2, 4, 2, -1, -6, -4, -5...
$ arr_time  <int> 830, 850, 923, 1004, 81...
$ sched_arr_time <int> 819, 830, 850, 1022, 83...
$ arr_delay <dbl> 11, 20, 33, -18, -25, 1...
$ carrier   <chr> "UA", "UA", "AA", "B6",...
$ flight    <int> 1545, 1714, 1141, 725, ...
$ tailnum   <chr> "N14228", "N24211", "N6...
$ origin    <chr> "EWR", "LGA", "JFK", "J...
$ dest      <chr> "IAH", "IAH", "MIA", "B...
$ air_time  <dbl> 227, 227, 160, 183, 116...
$ distance  <dbl> 1400, 1416, 1089, 1576,...
$ hour      <dbl> 5, 5, 5, 5, 6, 5, 6, 6,...
$ minute    <dbl> 15, 29, 40, 45, 0, 58, ...
$ time_hour <dtm> 2013-01-01 05:00:00, 2...
```

The dplyr package

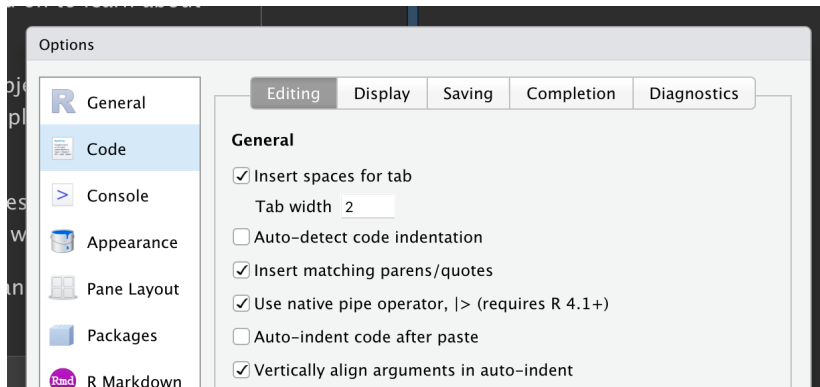
- The first argument is always a data frame.
- The subsequent arguments typically describe which columns to operate on, using the variable names (without quotes).
- The output is always a new data frame.
- Each verb operates on either
 - rows,
 - columns,
 - groups, or
 - tables

The pipe

$$\begin{aligned}x &|> f(y) &\Leftrightarrow& f(x, y) \\x &|> f(y) |> g(z) &\Leftrightarrow& g(f(x, y), z)\end{aligned}$$

```
flights |>
  filter(dest == "IAH") |>
  group_by(year, month, day) |>
  summarize(
    arr_delay = mean(arr_delay, na.rm = TRUE)
  )
```

Global options



- Enable “Use native pipe operator.”
- This will enable you to produce the pipe with `ctrl-shift-M`

filter

```
flights |>
  filter(dep_delay > 120)
#> # A tibble: 9,723 × 19
#>   year month   day dep_time sched_dep_time dep_delay arr
#>   <int> <int> <int>   <int>         <int>      <dbl>
#> 1  2013     1     1     848           1835      853
#> 2  2013     1     1     957           733      144
#> 3  2013     1     1    1114           900      134
#> 4  2013     1     1    1540          1338      122
#> 5  2013     1     1    1815          1325      290
#> 6  2013     1     1    1842          1422      260
#> # i 9,717 more rows
#> # i 11 more variables: arr_delay <dbl>, carrier <chr>, fl
```


arrange

```
flights |>
  arrange(year, month, day, dep_time)
#> # A tibble: 336,776 × 19
#>   year month   day dep_time sched_dep_time dep_delay arr
#>   <int> <int> <int>   <int>         <int>       <dbl> arr
#> 1  2013     1     1     517           515         2
#> 2  2013     1     1     533           529         4
#> 3  2013     1     1     542           540         2
#> 4  2013     1     1     544           545        -1
#> 5  2013     1     1     554           600        -6
#> 6  2013     1     1     554           558        -4
#> # i 9,717 more rows
#> # i 11 more variables: arr_delay <dbl>, carrier <chr>, fl
```

distinct

```
# Find all unique origin and destination pairs
flights |>
  distinct(origin, dest)
#> # A tibble: 224 × 2
#>   origin dest
#>   <chr>  <chr>
#> 1 EWR    IAH
#> 2 LGA    IAH
#> 3 JFK    MIA
#> 4 JFK    BQN
#> 5 LGA    ATL
#> 6 EWR    ORD
#> # i 218 more rows
```

count

```
flights |>
  count(origin, dest, sort = TRUE)
#> # A tibble: 224 × 3
#>   origin dest      n
#>   <chr>  <chr> <int>
#> 1 JFK    LAX    11262
#> 2 LGA    ATL    10263
#> 3 LGA    ORD     8857
#> 4 JFK    SFO     8204
#> 5 LGA    CLT     6168
#> 6 EWR    ORD     6100
#> # i 218 more rows
```

Do exercise 6

mutate

```
flights |>
  mutate(
    gain = dep_delay - arr_delay,
    speed = distance / air_time * 60,
    .before = 1
  )
#> # A tibble: 336,776 × 21
#>   gain speed  year month  day dep_time sched_dep_time c
#>   <dbl> <dbl> <int> <int> <int>   <int>           <int>
#> 1    -9  370.  2013     1     1     517             515
#> 2   -16  374.  2013     1     1     533             529
#> 3   -31  408.  2013     1     1     542             540
#> 4    17  517.  2013     1     1     544             545
#> 5    19  394.  2013     1     1     554             600
#> 6   -16  288.  2013     1     1     554             558
#> # i 336,770 more rows
#> # i 12 more variables: sched_arr_time<int>, arr_delay<int>
```

select

```
flights |>  
  select(year, month, day)
```

```
flights |>  
  select(year:day)
```

```
flights |>  
  select(!year:day)
```

```
flights |>  
  select(where(is.character))
```

- `starts_with("abc")`: matches names that begin with "abc".
- `ends_with("xyz")`: matches names that end with "xyz".
- `contains("ijk")`: matches names that contain "ijk".
- `num_range("x", 1:3)`: matches x1, x2 and x3.

rename

```
flights |>
  rename(tail_num = tailnum)
#> # A tibble: 336,776 × 19
#>   year month   day dep_time sched_dep_time dep_delay arr
#>   <int> <int> <int>   <int>         <int>         <dbl> arr
#> 1  2013     1     1     517           515           2   2
#> 2  2013     1     1     533           529           4   4
#> 3  2013     1     1     542           540           2   2
#> 4  2013     1     1     544           545          -1  -1
#> 5  2013     1     1     554           600          -6  -6
#> 6  2013     1     1     554           558          -4  -4
```

relocate

```
flights |>
  relocate(time_hour, air_time)
#> # A tibble: 336,776 × 19
#>   time_hour          air_time  year month   day dep_time
#>   <dtm>          <dbl> <int> <int> <int>   <int>
#> 1 2013-01-01 05:00:00     227  2013     1     1     517
#> 2 2013-01-01 05:00:00     227  2013     1     1     533
#> 3 2013-01-01 05:00:00     160  2013     1     1     542
#> 4 2013-01-01 05:00:00     183  2013     1     1     544
#> 5 2013-01-01 06:00:00     116  2013     1     1     554
#> 6 2013-01-01 05:00:00     150  2013     1     1     554
```


Do exercise 7

The Pipe

```
flights |>
  filter(dest == "IAH") |>
  mutate(speed = distance / air_time * 60) |>
  select(year:day, dep_time, carrier, flight, speed) |>
  arrange(desc(speed))
```

#> # A tibble: 7,198 × 7

#>	year	month	day	dep_time	carrier	flight	speed
#>	<int>	<int>	<int>	<int>	<chr>	<int>	<dbl>
#> 1	2013	7	9	707	UA	226	522.
#> 2	2013	8	27	1850	UA	1128	521.
#> 3	2013	8	28	902	UA	1711	519.
#> 4	2013	8	28	2122	UA	1022	519.
#> 5	2013	6	11	1628	UA	1178	515.
#> 6	2013	8	27	1017	UA	333	515.

The Pipe vs. function nesting

```
arrange(  
  select(  
    mutate(  
      filter(  
        flights,  
        dest == "IAH"  
      ),  
      speed = distance / air_time * 60  
    ),  
    year:day, dep_time, carrier, flight, speed  
  ),  
  desc(speed)  
)
```

The Pipe vs. assignment to temporaries

```
flights1 <- filter(flights, dest == "IAH")  
flights2 <- mutate(flights1, speed = distance / air_time * 60)  
flights3 <- select(flights2, year:day, dep_time, carrier, flight, speed)  
arrange(flights3, desc(speed))
```

group_by

```
flights |>
  group_by(month)
#> # A tibble: 336,776 × 19
#> # Groups:   month [12]
#>   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
#>   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
#> 1  2013     1     1     517           515           2     830           81
#> 2  2013     1     1     533           529           4     850           83
#> 3  2013     1     1     542           540           2     923           85
#> 4  2013     1     1     544           545          -1    1004          102
#> 5  2013     1     1     554           600          -6     812           83
#> 6  2013     1     1     554           558          -4     740           72
```

- All subsequent operations will now work “by month”

summarize

```
flights |>
  group_by(month) |>
  summarize(
    avg_delay = mean(dep_delay)
  )
```

#> # A tibble: 12 × 2

#>	month	avg_delay
#>	<int>	<dbl>
#> 1	1	NA
#> 2	2	NA
#> 3	3	NA
#> 4	4	NA
#> 5	5	NA
#> 6	6	NA

- We forgot `na.rm`

summarize

```
flights |>
  group_by(month) |>
  summarize(
    delay = mean(dep_delay, na.rm = TRUE)
  )
```

#> # A tibble: 12 × 2

#>	month	delay
#>	<int>	<dbl>
#> 1	1	10.0
#> 2	2	10.8
#> 3	3	13.2
#> 4	4	13.9
#> 5	5	13.0
#> 6	6	20.8

summarize with n

```
flights |>
  group_by(month) |>
  summarize(
    delay = mean(dep_delay, na.rm = TRUE),
    n = n()
  )
#> # A tibble: 12 × 3
#>   month delay      n
#>   <int> <dbl> <int>
#> 1     1  10.0 27004
#> 2     2  10.8 24951
#> 3     3  13.2 28834
#> 4     4  13.9 28330
#> 5     5  13.0 28796
#> 6     6  20.8 28243
```


slice

- `df |> slice_head(n = 1)`
takes the first row from each group.
- `df |> slice_tail(n = 1)`
takes the last row in each group.
- `df |> slice_min(x, n = 1)`
takes the row with the smallest value of column x.
- `df |> slice_max(x, n = 1)`
takes the row with the largest value of column x.
- `df |> slice_sample(n = 1)`
takes one random row.

slice

```
flights |>
  group_by(dest) |>
  slice_max(arr_delay, n = 1) |>
  relocate(dest)
```

#> # A tibble: 108 × 19

#> # Groups: dest [105]

#>	dest	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
#>	<chr>	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
#> 1	ABQ	2013	7	22	2145	2007	98	132
#> 2	ACK	2013	7	23	1139	800	219	1250
#> 3	ALB	2013	1	25	123	2000	323	229
#> 4	ANC	2013	8	17	1740	1625	75	2042
#> 5	ATL	2013	7	22	2257	759	898	121
#> 6	AUS	2013	7	10	2056	1505	351	2347

Groiuping by more than one variable


```
daily <- flights |>
  group_by(year, month, day)
daily
#> # A tibble: 336,776 × 19
#> # Groups:   year, month, day [365]
#>   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
#>   <int> <int> <int>   <int>         <int>         <dbl>    <int>         <int>
#> 1  2013     1     1     517           515           2      830           81
#> 2  2013     1     1     533           529           4      850           83
#> 3  2013     1     1     542           540           2      923           85
#> 4  2013     1     1     544           545          -1     1004          102
#> 5  2013     1     1     554           600          -6      812           83
#> 6  2013     1     1     554           558          -4      740           72
```

Case study

```
batters <- Lahman::Batting |>
  group_by(playerID) |>
  summarize(
    performance = sum(H, na.rm = TRUE) / sum(AB, na.rm = TRUE),
    n = sum(AB, na.rm = TRUE)
  )
batters
#> # A tibble: 20,166 × 3
#>   playerID performance      n
#>   <chr>         <dbl> <int>
#> 1 aardsda01      0         4
#> 2 aaronha01    0.305   12364
#> 3 aaronto01    0.229    944
#> 4 aasedo01      0         5
#> 5 abadan01     0.0952    21
#> 6 abadfe01     0.111     9
#> # i 20,160 more rows
```

Case study

```
batters |>  
  filter(n > 100) |>  
  ggplot(aes(x = n, y = performance)) +  
  geom_point(alpha = 1 / 10) +  
  geom_smooth(se = FALSE)
```



- The variation in performance is larger among players with fewer at-bats.
- There's a positive correlation between skill (performance) and opportunities to hit the ball (n) because teams want to give their best batters the most opportunities to hit the ball.

Finding the best batters

```
batters |>
  arrange(desc(performance))
#> # A tibble: 20,166 × 3
#>   playerID performance      n
#>   <chr>          <dbl> <int>
#> 1 abramge01             1     1
#> 2 alberan01             1     1
#> 3 banisje01             1     1
#> 4 bartocl01             1     1
#> 5 bassdo01              1     1
#> 6 birasst01             1     2
```

- http://varianceexplained.org/r/empirical_bayes_baseball/
- <https://www.evanmiller.org/how-not-to-sort-by-average-rating.html>

pivot_longer()

```
billboard
```

```
#> # A tibble: 317 × 79
```

```
#>   artist      track      date.entered   wk1   wk2   wk3   wk4   wk5  
#>   <chr>      <chr>      <date>      <dbl> <dbl> <dbl> <dbl> <dbl>  
#> 1 2 Pac      Baby Don't Cry (Ke... 2000-02-26      87    82    72    77    6  
#> 2 2Ge+her     The Hardest Part 0... 2000-09-02      91    87    92    NA    6  
#> 3 3 Doors Down Kryptonite      2000-04-08      81    70    68    67    6  
#> 4 3 Doors Down Loser      2000-10-21      76    76    72    69    6  
#> 5 504 Boyz     Wobble Wobble      2000-04-15      57    34    25    17    1  
#> 6 98~0         Give Me Just One N... 2000-08-19      51    39    34    26    6  
#> # i 311 more rows  
#> # i 71 more variables: wk6 <dbl>, wk7 <dbl>, wk8 <dbl>, wk9 <dbl>, ...
```

pivot_longer()

```
billboard |>
  pivot_longer(
    cols = starts_with("wk"),
    names_to = "week",
    values_to = "rank"
  )
```

#> # A tibble: 24,092 × 5

#>		artist	track		date.entered	week	rank
#>		<chr>	<chr>		<date>	<chr>	<dbl>
#>	1 2	Pac	Baby Don't Cry (Keep...		2000-02-26	wk1	87
#>	2 2	Pac	Baby Don't Cry (Keep...		2000-02-26	wk2	82
#>	3 2	Pac	Baby Don't Cry (Keep...		2000-02-26	wk3	72
#>	4 2	Pac	Baby Don't Cry (Keep...		2000-02-26	wk4	77
#>	5 2	Pac	Baby Don't Cry (Keep...		2000-02-26	wk5	87
#>	6 2	Pac	Baby Don't Cry (Keep...		2000-02-26	wk6	94
#>	7 2	Pac	Baby Don't Cry (Keep...		2000-02-26	wk7	99
#>	8 2	Pac	Baby Don't Cry (Keep...		2000-02-26	wk8	NA
#>	9 2	Pac	Baby Don't Cry (Keep...		2000-02-26	wk9	NA
#>	10 2	Pac	Baby Don't Cry (Keep...		2000-02-26	wk10	NA

#> # i 24,082 more rows

pivot_longer()

```
billboard |>
  pivot_longer(
    cols = starts_with("wk"),
    names_to = "week",
    values_to = "rank",
    values_drop_na = TRUE
  )
```

```
#> # A tibble: 5,307 × 5
```

#>	artist	track	date.entered	week	rank
#>	<chr>	<chr>	<date>	<chr>	<dbl>
#> 1	2 Pac	Baby Don't Cry (Keep...	2000-02-26	wk1	87
#> 2	2 Pac	Baby Don't Cry (Keep...	2000-02-26	wk2	82
#> 3	2 Pac	Baby Don't Cry (Keep...	2000-02-26	wk3	72
#> 4	2 Pac	Baby Don't Cry (Keep...	2000-02-26	wk4	77
#> 5	2 Pac	Baby Don't Cry (Keep...	2000-02-26	wk5	87
#> 6	2 Pac	Baby Don't Cry (Keep...	2000-02-26	wk6	94

```
#> # i 5,301 more rows
```

pivot_longer()

```
billboard_longer <- billboard |>
  pivot_longer(
    cols = starts_with("wk"),
    names_to = "week",
    values_to = "rank",
    values_drop_na = TRUE
  ) |>
  mutate(
    week = parse_number(week)
  )
```


```
billboard_longer
```

```
#> # A tibble: 5,307 × 5
```

```
#>   artist track      date.entered  week  rank
#>   <chr>  <chr>      <date>      <dbl> <dbl>
#> 1 2 Pac   Baby Don't Cry (Keep... 2000-02-26      1     87
#> 2 2 Pac   Baby Don't Cry (Keep... 2000-02-26      2     82
#> 3 2 Pac   Baby Don't Cry (Keep... 2000-02-26      3     72
#> 4 2 Pac   Baby Don't Cry (Keep... 2000-02-26      4     77
#> 5 2 Pac   Baby Don't Cry (Keep... 2000-02-26      5     87
#> 6 2 Pac   Baby Don't Cry (Keep... 2000-02-26      6     94
#> # i 5,301 more rows
```

pivot_longer()

```
billboard_longer |>  
  ggplot(aes(x = week, y = rank, group = track)) +  
  geom_line(alpha = 0.25) +  
  scale_y_reverse()
```




pivot_longer()

```
df <- tribble(
  ~id, ~bp1, ~bp2,
  "A", 100, 120,
  "B", 140, 115,
  "C", 120, 125
)
df |>
  pivot_longer(
    cols = bp1:bp2,
    names_to = "measurement",
    values_to = "value"
  )
#> # A tibble: 6 × 3
#>   id      measurement value
#>   <chr> <chr>         <dbl>
#> 1 A      bp1           100
#> 2 A      bp2           120
#> 3 B      bp1           140
#> 4 B      bp2           115
#> 5 C      bp1           120
#> 6 C      bp2           125
```

pivot_longer()

```
df |>  
  pivot_longer(  
    cols = bp1:bp2,  
    names_to = "measurement",  
    values_to = "value"  
  )
```

id	bp1	bp2
A	100	120
B	140	115
C	120	125




id	name	value
A	bp1	100
A	bp2	120
B	bp1	140
B	bp2	115
C	bp1	120
C	bp2	125

pivot_longer()

```
df |>  
  pivot_longer(  
    cols = bp1:bp2,  
    names_to = "measurement",  
    values_to = "value"  
  )
```

id	bp1	bp2
A	100	120
B	140	115
C	120	125




id	name	value
A	bp1	100
A	bp2	120
B	bp1	140
B	bp2	115
C	bp1	120
C	bp2	125

pivot_longer()

```
df |>  
  pivot_longer(  
    cols = bp1:bp2,  
    names_to = "measurement",  
    values_to = "value"  
  )
```

id	bp1	bp2
A	100	120
B	140	115
C	120	125



id	name	value
A	bp1	100
A	bp2	120
B	bp1	140
B	bp2	115
C	bp1	120
C	bp2	125